

# CS294-112 Deep Reinforcement Learning HW2: Policy Gradients due September 19th 2018, 11:59 pm

## 1 Introduction

The goal of this assignment is to experiment with policy gradient and its variants, including variance reduction methods. Your goals will be to set up policy gradient for both continuous and discrete environments and experiment with variance reduction tricks, including implementing reward-to-go and neural network baselines.

Turn in your report and code for the full homework as described in Problem 2 by September 19, 2018.

## 2 Review

Recall that the reinforcement learning objective is to learn a  $\theta^*$  that maximizes the objective function:

$$J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}(\tau)} [r(\tau)] \tag{1}$$

where

$$\pi_{\theta}(\tau) = p(s_1, a_1, \dots, s_T, a_T) = p(s_1) \pi_{\theta}(a_1 | s_1) \prod_{t=2}^T p(s_t | s_{t-1}, a_{t-1}) \pi_{\theta}(a_t | s_t)$$

and

$$r(\tau) = r(s_1, a_1, \dots, s_T, a_T) = \sum_{t=1}^T r(s_t, a_t).$$

The policy gradient approach is to directly take the gradient of this objective:

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \int \pi_{\theta}(\tau) r(\tau) d\tau \quad (2)$$

$$= \int \pi_{\theta}(\tau) \nabla_{\theta} \log \pi_{\theta}(\tau) r(\tau) d\tau. \quad (3)$$

In practice, the expectation over trajectories  $\tau$  can be approximated from a batch of  $N$  sampled trajectories:

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \log \pi_{\theta}(\tau_i) r(\tau_i) \quad (4)$$

$$= \frac{1}{N} \sum_{i=1}^N \left( \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_{it}|s_{it}) \right) \left( \sum_{t=1}^T r(s_{it}, a_{it}) \right). \quad (5)$$

Here we see that the policy  $\pi_{\theta}$  is a probability distribution over the action space, conditioned on the state. In the agent-environment loop, the agent samples an action  $a_t$  from  $\pi_{\theta}(\cdot|s_t)$  and the environment responds with a reward  $r(s_t, a_t)$ .

One way to reduce the variance of the policy gradient is to exploit causality: the notion that the policy cannot affect rewards in the past, yielding following the modified objective, where the sum of rewards here is a sample estimate of the  $Q$  function, known as the “reward-to-go:”

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_{it}|s_{it}) \left( \sum_{t'=t}^T r(s_{it'}, a_{it'}) \right). \quad (6)$$

Multiplying a discount factor  $\gamma$  to the rewards can be interpreted as encouraging the agent to focus on rewards closer in the future, which can also be thought of as a means for reducing variance (because there is more variance possible futures further into the future). We saw in lecture that the discount factor can be incorporated in two ways.

The first way applies the discount on the rewards from full trajectory:

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left( \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_{it}|s_{it}) \right) \left( \sum_{t=1}^T \gamma^{t-1} r(s_{it}, a_{it}) \right) \quad (7)$$

and the second way applies the discount on the “reward-to-go:”

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_{it}|s_{it}) \left( \sum_{t'=t}^T \gamma^{t'-t} r(s_{it'}, a_{it'}) \right). \quad (8)$$

We have seen in lecture that subtracting a baseline that is a constant with respect to  $\tau$  from the sum of rewards

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \mathbb{E}_{\tau \sim \pi_{\theta}(\tau)} [r(\tau) - b] \quad (9)$$

leaves the policy gradient unbiased because

$$\nabla_{\theta} \mathbb{E}_{\tau \sim \pi_{\theta}(\tau)} [b] = 0.$$

In this assignment, we will implement a value function  $V_{\phi}^{\pi}$  which acts as a *state-dependent* baseline. The value function is trained to approximate the sum of future rewards starting from a particular state:

$$V_{\phi}^{\pi}(s_t) \approx \sum_{t'=t}^T \mathbb{E}_{\pi_{\theta}} [r(s_{t'}, a_{t'}) | s_t], \quad (10)$$

so the approximate policy gradient now looks like this:

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_{it} | s_{it}) \left( \left( \sum_{t'=t}^T \gamma^{t'-t} r(s_{it'}, a_{it'}) \right) - V_{\phi}^{\pi}(s_{it}) \right). \quad (11)$$

**Problem 1. State-dependent baseline:** In lecture we saw that the policy gradient is unbiased if the baseline is a constant with respect to  $\tau$  (Equation 9). The purpose of this problem is to help convince ourselves that subtracting a state-dependent baseline from the return keeps the policy gradient unbiased. Using the [law of iterated expectations](#) show that the policy gradient is still unbiased if the baseline  $b$  is function of a state at a particular timestep of  $\tau$  (Equation 11). Please answer the questions below in L<sup>A</sup>T<sub>E</sub>X in your report.

(a) Note that by [linearity of expectation](#) the objective can be written as:

$$\begin{aligned} J(\theta) &= \mathbb{E}_{\tau \sim \pi_{\theta}(\tau)} [r(\tau)] \\ &= \mathbb{E}_{\tau \sim \pi_{\theta}(\tau)} \left[ \sum_{i=1}^T r(s_i, a_i) \right] \\ &= \sum_{i=1}^T \mathbb{E}_{(s_i, a_i) \sim p(s_i, a_i)} [r(s_i, a_i)] \end{aligned}$$

when we subtract the baseline  $b(s_t)$ , the objective becomes:

$$= \sum_{i=1}^T \mathbb{E}_{(s_i, a_i) \sim p(s_i, a_i)} [r(s_i, a_i) - b(s_i)].$$

Please show that

$$\nabla_{\theta} \sum_{i=1}^T \mathbb{E}_{(s_i, a_i) \sim p(s_i, a_i)} [b(s_i)] = 0.$$

- (b) An alternative approach is to look at the entire trajectory and consider a particular timestep  $t^* \in [1, T - 1]$  (the timestep  $T$  case would be very similar to part (a)).
- (a) We can exploit the conditional independency structure of  $\pi_\theta(\tau) = p(s_1, a_1, \dots, s_T, a_T)$  and use the law of iterated expectations to break Equation 12 into two expectations, where the outer expectation is over  $(s_1, a_1, \dots, a_{t^*} - 1, s_{t^*})$ , and the inner expectation is over the rest of the trajectory, conditioned on  $(s_1, a_1, \dots, a_{t^*} - 1, s_{t^*})$ . Explain why, for the inner expectation, conditioning on  $(s_1, a_1, \dots, a_{t^*} - 1, s_{t^*})$  is equivalent to conditioning only on  $s_{t^*}$ .
- (b) Using the iterated expectation described above, show that

$$\nabla_\theta \mathbb{E}_{\tau \sim \pi_\theta(\tau)} [b(s_{t^*})] = 0. \quad (12)$$

Since the policy gradient with respect to  $\theta$  can be decoupled as a summation of terms over timesteps  $t \in [1, T]$ , because we have shown that the policy gradient is unbiased for each of these terms, the entire policy gradient is also unbiased with respect to a vector of state-dependent baselines over the timesteps:  $[b(s_1), b(s_2), \dots, b(s_T)]$ .

## 3 Code Setup

### 3.1 Files

The starter code is available [here](#). The only file you need to modify in this homework is `train_pg_f18.py`. The files `logz.py` and `plots.py` are utility files; while you should look at them to understand their functionality, you will not modify them. For the Lunar Lander task, use the provided `lunar_lander.py` file instead of `gym/envs/box2d/lunar_lander.py`. After you fill in the appropriate methods, you should be able to just run `python train_pg_f18.py` with some command line options to perform the experiments. To visualize the results, you can run `python plot.py path/to/logdir`.

### 3.2 Overview

The function `train_PG` is used to perform the actual training for policy gradient. The parameters passed into this function specify the algorithm's hyperparameters and environment. The `Agent` class contains methods that define the computation graph, sample trajectories, estimate returns, and update the parameters of the policy.

At a high level, the dataflow of the code is structured like this:

1. *Build a static computation graph* in Tensorflow.

2. *Set up a Tensorflow session* to initialize the parameters of the computation graph. This is the only time you need to set up the session.

Then we will repeat Steps 3 through 5 for  $N$  iterations:

3. *Sample trajectories* by executing the Tensorflow op that samples an action given an observation from the environment. Collect the states, actions, and rewards as numpy variables.
4. *Estimate returns* in numpy (estimated Q values, baseline predictions, advantages).
5. *Update parameters* by executing the Tensorflow op that updates the parameters given what you computed in Step 4.

## 4 Constructing the computation graph

**Problem 2. Neural networks:** We will now begin to implement a neural network that parametrizes  $\pi_\theta$ .

- (a) Implement the utility function, `build_mlp`, which will build a feedforward neural network with fully connected units (Hint: use `tf.layers.dense`). Test it to make sure that it produces outputs of the expected size and shape. **You do not need to include anything in your write-up about this**, it will just make your life easier.
- (b) Next, implement the method `Agent.build_computation_graph`. At this point, you only need to implement the parts with the “Problem 2” header.
  - (i) Define the placeholder for the advantages in `Agent.define_placeholders`. We have already defined placeholders for the observations and actions. The advantages will correspond to  $r(\tau)$  in the policy gradient, which may or may not include a subtracted baseline value.
  - (ii) Create the symbolic operation `Agent.policy_forward_pass`: This outputs the parameters of a distribution  $\pi_\theta(a|s)$ . In this homework, when the distribution is over discrete actions these parameters will be the logits of a categorical distribution, and when the distribution is over continuous actions these parameters will be the mean and the log standard deviation of a multivariate Gaussian distribution. This operation will be an input to `Agent.sample_action` and `Agent.get_log_prob`.
  - (iii) Create the symbolic operation `Agent.sample_action`: This produces a Tensorflow op, `self.sy_sampled_ac` that samples an action from  $\pi_\theta(a|s)$ . This operation will be called in `Agent.sample_trajectories`.
  - (iv) Create the symbolic operation `Agent.get_log_prob`: Given an action that the agent took in the environment, this computes the log probability of that action

under  $\pi_\theta(a|s)$ . This will be used in the loss function.

- (v) In `Agent.build_computation_graph` implement a loss function (which call use the result from `Agent.get_log_prob`) to whose gradient is

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \nabla_\theta \log \pi_\theta(\tau_i) r(\tau_i).$$

## 5 Implement Policy Gradient

### 5.1 Implementing the policy gradient loop

**Problem 3. Policy Gradient:** Recall from lecture that an RL algorithm can viewed as consisting of three parts, which are reflected in the training loop of `train_PG`:

1. `Agent.sample_trajectories`: Generate samples (e.g. run the policy to collect trajectories consisting of state transitions  $(s, a, s', r)$ )
2. `Agent.estimate_return`: Estimate the return (e.g. sum together discounted rewards from the trajectories, or learn a model that predicts expected total future discounted reward)
3. `Agent.update_parameters`: Improve the policy (e.g. update the parameters of the policy with policy gradient)

In our implementation, for clarity we will update the parameters of the value function baseline also in the third step (`Agent.update_parameters`), rather than in the second step (as was described in lecture). You only need to implement the parts with the “Problem 3” header.

- (a) **Sample trajectories:** In `Agent.sample_trajectories`, use the Tensorflow session to call `self.sy_sampled_ac` to sample an action given an observation from the environment.
- (b) **Estimate return:** We will now implement  $r(\tau)$  from Equation 1. Please implement the method `Agent.sum_of_rewards`, which will return a sample estimate of the discounted return, for both the full-trajectory (Equation 7) case, where

$$r(\tau_i) = \sum_{t=1}^T \gamma^{t-1} r(s_{it}, a_{it})$$

and for the “reward-to-go” case (Equation 8) where

$$r(\tau_i) = \sum_{t'=t}^T \gamma^{t'-t} r(s_{it'}, a_{it'}).$$

In `Agent.estimate_return`, normalize the advantages to have a mean of zero and a standard deviation of one. This is a trick for reducing variance.

- (c) **Update parameters:** In `Agent.update_parameters` use the Tensorflow session to call the update operation `self.update_op` to update the parameters of the policy. You will need to figure out the inputs to `feed_dict`.

## 5.2 Experiments

After you have implemented the code, we will run experiments to get a feel for how different settings impact the performance of policy gradient methods.

**Problem 4. CartPole:** Run the PG algorithm in the discrete `CartPole-v0` environment from the command line as follows:

```
python train_pg_fl18.py CartPole-v0 -n 100 -b 1000 -e 3 -dna --exp_name
sb_no_rtg_dna
python train_pg_fl18.py CartPole-v0 -n 100 -b 1000 -e 3 -rtg -dna --exp_name
sb_rtg_dna
python train_pg_fl18.py CartPole-v0 -n 100 -b 1000 -e 3 -rtg --exp_name
sb_rtg_na
python train_pg_fl18.py CartPole-v0 -n 100 -b 5000 -e 3 -dna --exp_name
lb_no_rtg_dna
python train_pg_fl18.py CartPole-v0 -n 100 -b 5000 -e 3 -rtg -dna --exp_name
lb_rtg_dna
python train_pg_fl18.py CartPole-v0 -n 100 -b 5000 -e 3 -rtg --exp_name
lb_rtg_na
```

What's happening there:

- `-n` : Number of iterations.
- `-b` : Batch size (number of state-action pairs sampled while acting according to the current policy at each iteration).
- `-e` : Number of experiments to run with the same configuration. Each experiment will start with a different randomly initialized policy, and have a different stream of random numbers.
- `-dna` : Flag: if present, sets `normalize_advantages` to `False`. Otherwise, by default, `normalize_advantages=True`.
- `-rtg` : Flag: if present, sets `reward_to_go=True`. Otherwise, by default, `reward_to_go=False`.
- `--exp_name` : Name for experiment, which goes into the name for the data directory.

Various other command line arguments will allow you to set batch size, learning rate, network architecture (number of hidden layers and the size of the hidden layers—for `CartPole`, you can use one hidden layer with 32 units), and more.

### Deliverables for report:

- Graph the results of your experiments **using the plot.py file we provide**. Create two graphs.
  - In the first graph, compare the learning curves (average return at each iteration) for the experiments prefixed with `sb_`. (The small batch experiments.)
  - In the second graph, compare the learning curves for the experiments prefixed with `lb_`. (The large batch experiments.)
- Answer the following questions briefly:
  - Which gradient estimator has better performance without advantage-centering—the trajectory-centric one, or the one using reward-to-go?
  - Did advantage centering help?
  - Did the batch size make an impact?
- Provide the exact command line configurations you used to run your experiments. (To verify batch size, learning rate, architecture, and so on.)

### What to Expect:

- The best configuration of CartPole in both the large and small batch cases converge to a maximum score of 200.

**Problem 5. InvertedPendulum:** Run experiments in InvertedPendulum-v2 continuous control environment as follows:

```
python train_pg_fl8.py InvertedPendulum-v2 -ep 1000 --discount 0.9 -n 100 -e 3  
-l 2 -s 64 -b <b*> -lr <r*> -rtg --exp_name hc_b<b*>_r<r*>
```

where your task is to find the smallest batch size  $b^*$  and largest learning rate  $r^*$  that gets to optimum (maximum score of 1000) in less than 100 iterations. The policy performance may fluctuate around 1000 – this is fine. The precision of  $b^*$  and  $r^*$  need only be one significant digit.

### Deliverables:

- Given the  $b^*$  and  $r^*$  you found, provide a learning curve where the policy gets to optimum (maximum score of 1000) in less than 100 iterations. (This may be for a single random seed, or averaged over multiple.)
- Provide the exact command line configurations you used to run your experiments.



## 6 Implement Neural Network Baselines

For the rest of the assignment we will use “reward-to-go.”

**Problem 6. Neural network baseline:** We will now implement a value function as a state-dependent neural network baseline. The sections in the code are marked by “Problem 6.”

- (a) In `Agent.build_computation_graph` implement  $V_{\phi}^{\pi}$ , a neural network that predicts the expected return conditioned on a state. Also implement the loss function to train this network and its update operation `self.baseline_op`.
- (b) In `Agent.compute_advantage`, use the neural network to predict the expected state-conditioned return (use the session to call `self.baseline_prediction`), normalize it to match the statistics of the current batch of “reward-to-go”, and subtract this value from the “reward-to-go” to yield an estimate of the advantage. This implements  $\left(\sum_{t'=t}^T \gamma^{t'-t} r(s_{it'}, a_{it'})\right) - V_{\phi}^{\pi}(s_{it})$ .
- (c) In `Agent.update_parameters`, update the parameters of the the neural network baseline by using the Tensorflow session to call `self.baseline_op`. “Rescale” the target values for the neural network baseline to have a mean of zero and a standard deviation of one.

## 7 More Complex Tasks

**Note:** The following tasks would take quite a bit of time to train. Please start early!

**Problem 7: LunarLander** For this problem, you will use your policy gradient implementation to solve `LunarLanderContinuous-v2`. Use an episode length of 1000. The purpose of this problem is to help you debug your baseline implementation. Run the following command:

```
python train_pg_fl8.py LunarLanderContinuous-v2 -ep 1000 --discount 0.99 -n
100 -e 3 -l 2 -s 64 -b 40000 -lr 0.005 -rtg --nn_baseline --exp_name
ll_b40000_r0.005
```

**Deliverables:**

- Plot a learning curve for the above command. You should expect to achieve an average return of around 180.

**Problem 8: HalfCheetah** For this problem, you will use your policy gradient implementation to solve `HalfCheetah-v2`. Use an episode length of 150, which is shorter than the default of 1000 for `HalfCheetah` (which would speed up your training significantly). Search

over batch sizes  $b \in [10000, 30000, 50000]$  and learning rates  $r \in [0.005, 0.01, 0.02]$  to replace  $\langle b \rangle$  and  $\langle r \rangle$  below:

```
python train_pg_f18.py HalfCheetah-v2 -ep 150 --discount 0.9 -n 100 -e 3 -l 2
-s 32 -b  $\langle b \rangle$  -lr  $\langle r \rangle$  --exp_name hc_b $\langle b \rangle$ _r $\langle r \rangle$ 
```

### Deliverables:

- How did the batch size and learning rate affect the performance?
- Once you've found suitable values of  $b$  and  $r$  among those choices (let's call them  $b^*$  and  $r^*$ ), use  $b^*$  and  $r^*$  and run the following commands (remember to replace the terms in the angle brackets):

```
python train_pg_f18.py HalfCheetah-v2 -ep 150 --discount 0.9 -n 100 -e 3
-l 2 -s 32 -b  $\langle b^* \rangle$  -lr  $\langle r^* \rangle$  --exp_name hc_b $\langle b^* \rangle$ _r $\langle r^* \rangle$ 
python train_pg_f18.py HalfCheetah-v2 -ep 150 --discount 0.9 -n 100 -e 3
-l 2 -s 32 -b  $\langle b^* \rangle$  -lr  $\langle r^* \rangle$  -rtg --exp_name hc_b $\langle b^* \rangle$ _r $\langle r^* \rangle$ 
python train_pg_f18.py HalfCheetah-v2 -ep 150 --discount 0.9 -n 100 -e 3
-l 2 -s 32 -b  $\langle b^* \rangle$  -lr  $\langle r^* \rangle$  --nn_baseline --exp_name hc_b $\langle b^* \rangle$ _r $\langle r^* \rangle$ 
python train_pg_f18.py HalfCheetah-v2 -ep 150 --discount 0.9 -n 100 -e 3
-l 2 -s 32 -b  $\langle b^* \rangle$  -lr  $\langle r^* \rangle$  -rtg --nn_baseline --exp_name hc_b $\langle b^* \rangle$ _r $\langle r^* \rangle$ 
*>
```

The run with reward-to-go and the baseline should achieve an average score close to 200. Provide a single plot plotting the learning curves for all four runs.

## 8 Bonus!

Choose any (or all) of the following:

- A serious bottleneck in the learning, for more complex environments, is the sample collection time. In `train_pg_f18.py`, we only collect trajectories in a single thread, but this process can be fully parallelized across threads to get a useful speedup. Implement the parallelization and report on the difference in training time.
- Implement GAE- $\lambda$  for advantage estimation.<sup>1</sup> Run experiments in a MuJoCo gym environment to explore whether this speeds up training. (Walker2d-v1 may be good for this.)
- In PG, we collect a batch of data, estimate a single gradient, and then discard the data and move on. Can we potentially accelerate PG by taking multiple gradient descent steps with the same batch of data? Explore this option and report on your results. Set up a fair comparison between single-step PG and multi-step PG on at least one MuJoCo gym environment.

---

<sup>1</sup><https://arxiv.org/abs/1506.02438>

## 9 Submission

Your report should be a document containing

- (a) Your mathematical response (written in  $\text{\LaTeX}$ ) for Problem 1.
- (b) All graphs requested in Problems 4, 5, 7, and 8.
- (c) Answers to short explanation questions in section 5 and 7.
- (d) All command-line expressions you used to run your experiments.
- (e) (Optionally) Your bonus results (command-line expressions, graphs, and a few sentences that comment on your findings).

Please also submit your modified `train_pg_f18.py` file. If your code includes additional files, provide a zip file including your `train_pg_f18.py` and all other files needed to run your code. Please include a `README.md` with instructions needed to exactly duplicate your results (including command-line expressions).

Turn in your assignment by September 19th 11:59pm on Gradescope. Upload the zip file with your code to **HW2 Code**, and upload the PDF of your report to **HW2**.